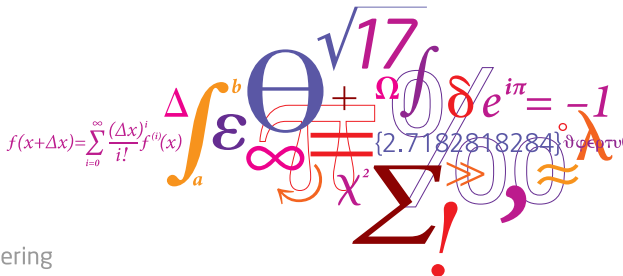


Variational inference

Rico Krueger

Filipe Rodrigues



Learning objectives

- Explain the concept of deterministic approximate inference.
- Explain variational inference.
- Describe differences between variational inference and expectation propagation.
- Derive and explain the evidence lower bound.
- Apply mean-field variational inference.
- Explain and apply Variational Bayesian Expectation-Maximization.

Outline



- Introduction
- Variational inference
- Variational Bayesian Expectation-Maximization

Approximate inference

- Suppose we wish to compute the posterior distribution of the latent variables \mathbf{z} given some observed data \mathbf{x} : $p(\mathbf{z}|\mathbf{x})$
- For many problems of interest exact posterior inference is **intractable**

$$\begin{aligned}
 \text{posterior} \quad \overbrace{p(\mathbf{z}|\mathbf{x})} &= \frac{\overbrace{p(\mathbf{x}, \mathbf{z})}^{\text{joint}}}{\underbrace{p(\mathbf{x})}_{\text{evidence}}} = \frac{\overbrace{p(\mathbf{x}|\mathbf{z})}^{\text{likelihood}} \overbrace{p(\mathbf{z})}^{\text{prior}}}{\underbrace{p(\mathbf{x})}_{\text{evidence}}} = \frac{\overbrace{p(\mathbf{x}|\mathbf{z})}^{\text{likelihood}} \overbrace{p(\mathbf{z})}^{\text{prior}}}{\underbrace{\sum_{\mathbf{z}} p(\mathbf{x}|\mathbf{z}) p(\mathbf{z})}_{\text{evidence}}}
 \end{aligned}$$

- Cannot determine the posterior distribution analytically
- Cannot even compute expectations with respect to the posterior

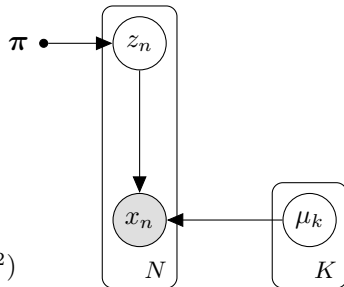
Approximate inference

- Suppose we wish to compute the posterior distribution of the latent variables \mathbf{z} given some observed data \mathbf{x} : $p(\mathbf{z}|\mathbf{x})$
- For many problems of interest exact posterior inference is **intractable**

$$\underbrace{p(\mathbf{z}|\mathbf{x})}_{\text{posterior}} = \frac{\overbrace{p(\mathbf{x}, \mathbf{z})}^{\text{joint}}}{\underbrace{p(\mathbf{x})}_{\text{evidence}}} = \frac{\overbrace{p(\mathbf{x}|\mathbf{z})}^{\text{likelihood}} \overbrace{p(\mathbf{z})}^{\text{prior}}}{\underbrace{p(\mathbf{x})}_{\text{evidence}}} = \frac{\overbrace{p(\mathbf{x}|\mathbf{z})}^{\text{likelihood}} \overbrace{p(\mathbf{z})}^{\text{prior}}}{\underbrace{\sum_{\mathbf{z}} p(\mathbf{x}|\mathbf{z}) p(\mathbf{z})}_{\text{evidence}}}$$

- Cannot determine the posterior distribution analytically
- Cannot even compute expectations with respect to the posterior
- Reasons for intractability:
 - Dimensionality of the latent space is too high to work with directly
 - Posterior distribution has a highly complex form for which expectations are not analytically tractable
- We must resort to **approximate inference** methods!

A practical example: Bayesian Gaussian mixture model



- For the sake of simplicity, assume π is given (fixed)

1 For each cluster k :

a) Draw cluster center $\mu_k \sim \mathcal{N}(\mu_k | 0, \tau^2)$

2 For each data point $1, \dots, N$:

a) Draw cluster assignment $z_n \sim \text{Multinomial}(z_n | \pi)$

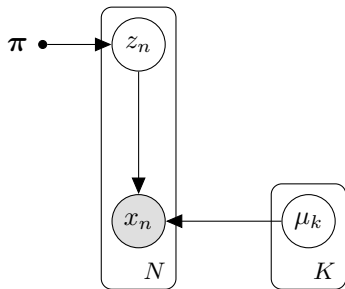
b) Draw observed data point $x_n \sim \mathcal{N}(x_n | \mu_{z_n}, \sigma^2)$

Note

For simplicity, we assumed the cluster variances σ^2 to be fixed.

A practical example: Bayesian Gaussian mixture model

- Joint distribution is given by (ignoring the fixed parameters)

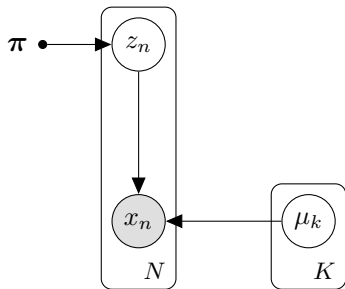


$$p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x}) = \left(\prod_{k=1}^K p(\mu_k) \right) \prod_{n=1}^N p(z_n | \boldsymbol{\pi}) p(x_n | z_n, \boldsymbol{\mu})$$

where $\boldsymbol{\mu} = \{\mu_1, \dots, \mu_K\}$, $\mathbf{x} = \{x_1, \dots, x_N\}$ and $\mathbf{z} = \{z_1, \dots, z_N\}$

A practical example: Bayesian Gaussian mixture model

- Joint distribution is given by (ignoring the fixed parameters)



$$p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x}) = \left(\prod_{k=1}^K p(\mu_k) \right) \prod_{n=1}^N p(z_n | \boldsymbol{\pi}) p(x_n | z_n, \boldsymbol{\mu})$$

where $\boldsymbol{\mu} = \{\mu_1, \dots, \mu_K\}$, $\mathbf{x} = \{x_1, \dots, x_N\}$ and $\mathbf{z} = \{z_1, \dots, z_N\}$

- Our goal is to compute the posterior over $\boldsymbol{\mu}$ and \mathbf{z}

$$p(\boldsymbol{\mu}, \mathbf{z} | \mathbf{x}) = \frac{p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}{\int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}$$

A practical example: Bayesian Gaussian mixture model

$$p(\boldsymbol{\mu}, \mathbf{z} | \mathbf{x}) = \frac{p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}{\int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}$$

- The numerator is easy to evaluate for any configuration of the hidden variables

A practical example: Bayesian Gaussian mixture model

$$p(\boldsymbol{\mu}, \mathbf{z}|\mathbf{x}) = \frac{p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}{\int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}$$

- The numerator is easy to evaluate for any configuration of the hidden variables
- The problem is the marginal likelihood in the denominator

$$\begin{aligned} p(\mathbf{x}) &= \int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x}) \\ &= \int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} \left(\prod_{k=1}^K p(\mu_k) \right) \prod_{n=1}^N p(z_n | \boldsymbol{\pi}) p(x_n | z_n, \boldsymbol{\mu}) \end{aligned}$$

A practical example: Bayesian Gaussian mixture model

$$p(\boldsymbol{\mu}, \mathbf{z}|\mathbf{x}) = \frac{p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}{\int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})}$$

- The numerator is easy to evaluate for any configuration of the hidden variables
- The problem is the marginal likelihood in the denominator

$$\begin{aligned} p(\mathbf{x}) &= \int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x}) \\ &= \int_{\boldsymbol{\mu}} \sum_{\mathbf{z}} \left(\prod_{k=1}^K p(\mu_k) \right) \prod_{n=1}^N p(z_n|\boldsymbol{\pi}) p(x_n|z_n, \boldsymbol{\mu}) \end{aligned}$$

- Taking advantage of the conditional independence of the z_n 's:

$$p(\mathbf{x}) = \int_{\boldsymbol{\mu}} \left(\prod_{k=1}^K p(\mu_k) \right) \prod_{n=1}^N \sum_{z_n} p(z_n|\boldsymbol{\pi}) p(x_n|z_n, \boldsymbol{\mu})$$

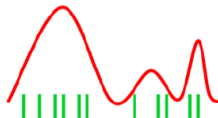
- But this integral is still **intractable** to compute!

Approximate inference: stochastic methods

- Obtain a set of **samples** $\mathbf{z}^{(s)}$, for $s \in \{1, \dots, S\}$, drawn independently from the distribution $p(\mathbf{z}|\mathbf{x})$



True distribution



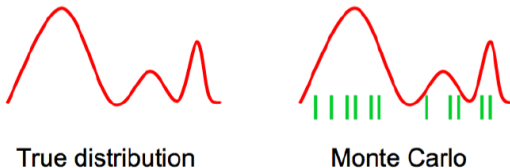
Monte Carlo

- Allows to approximate expectations as finite sums

$$\mathbb{E}_{p(\mathbf{z}|\mathbf{x})}[f(\mathbf{z})] \approx \frac{1}{S} \sum_{s=1}^S f(\mathbf{z}^{(s)})$$

Approximate inference: stochastic methods

- Obtain a set of **samples** $\mathbf{z}^{(s)}$, for $s \in \{1, \dots, S\}$, drawn independently from the distribution $p(\mathbf{z}|\mathbf{x})$



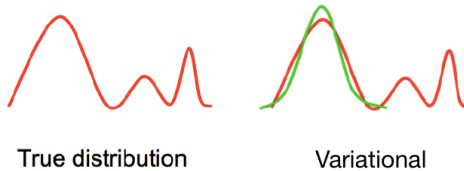
- Allows to approximate expectations as finite sums

$$\mathbb{E}_{p(\mathbf{z}|\mathbf{x})}[f(\mathbf{z})] \approx \frac{1}{S} \sum_{s=1}^S f(\mathbf{z}^{(s)})$$

- In the limit of infinite computational resources they can generate exact results
- Computationally demanding** and hard to scale to large datasets
- Many practical problems: determining convergence, number of samples, burn-in size, thinning, hard to diagnose, etc.

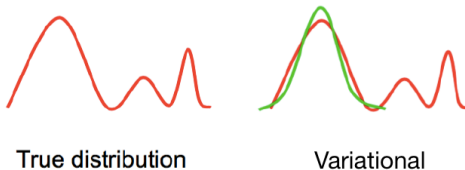
Approximate inference: variational methods

- Consider a family of tractable distributions $q(\mathbf{z}|\boldsymbol{\nu})$ - the **variational distribution**
- Find the **variational parameters** $\boldsymbol{\nu}$ that make $q(\mathbf{z}|\boldsymbol{\nu})$ as close as possible to the true posterior $p(\mathbf{z}|\mathbf{x})$



Approximate inference: variational methods

- Consider a family of tractable distributions $q(\mathbf{z}|\boldsymbol{\nu})$ - the **variational distribution**
- Find the **variational parameters** $\boldsymbol{\nu}$ that make $q(\mathbf{z}|\boldsymbol{\nu})$ as close as possible to the true posterior $p(\mathbf{z}|\mathbf{x})$



- Turns the inference problem into an **optimization** problem!
- Use $q(\mathbf{z}|\boldsymbol{\nu})$ with the fitted parameters as a proxy for the true posterior
 - To make predictions about future data
 - To investigate the posterior distribution of the hidden variables

Mean-field variational inference

- How should we choose a **tractable** family of distributions for $q(\mathbf{z})$?

Mean-field variational inference

- How should we choose a **tractable** family of distributions for $q(\mathbf{z})$?
 - Relax some constraints in the true distribution (e.g. dependencies)
 - For example, we can assume a **fully factorized** approximation

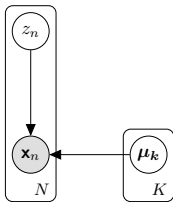
$$q(\mathbf{z}) = \prod_{m=1}^M q(z_m)$$

Mean-field variational inference

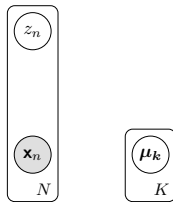
- How should we choose a **tractable** family of distributions for $q(\mathbf{z})$?
 - Relax some constraints in the true distribution (e.g. dependencies)
 - For example, we can assume a **fully factorized** approximation

$$q(\mathbf{z}) = \prod_{m=1}^M q(z_m)$$

- This is called a **mean-field approximation**



True joint distribution



Fully factorized approximation

- We can also assume that the distribution factorizes in groups of latent variables

Kullback-Leibler (KL) divergence

- What criteria defines the closeness between the two distributions?
 - **Kullback-Leibler (KL) divergence**

$$\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) = \int_{\mathbf{z}} q(\mathbf{z}) \log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} = \mathbb{E}_q \left[\log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right]$$

Kullback-Leibler (KL) divergence

- What criteria defines the closeness between the two distributions?
 - **Kullback-Leibler (KL) divergence**

$$\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) = \int_{\mathbf{z}} q(\mathbf{z}) \log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} = \mathbb{E}_q \left[\log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right]$$

- Notice that the KL divergence is an asymmetric measure

$$\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) \neq \mathbb{KL}(p(\mathbf{z}|\mathbf{x})||q(\mathbf{z}))$$

Kullback-Leibler (KL) divergence

- What criteria defines the closeness between the two distributions?
 - **Kullback-Leibler (KL) divergence**

$$\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) = \int_{\mathbf{z}} q(\mathbf{z}) \log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} = \mathbb{E}_q \left[\log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right]$$

- Notice that the KL divergence is an asymmetric measure

$$\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) \neq \mathbb{KL}(p(\mathbf{z}|\mathbf{x})||q(\mathbf{z}))$$

- The reverse KL, $\mathbb{KL}(p(\mathbf{z}|\mathbf{x})||q(\mathbf{z}))$, gives rise to a different approximation inference algorithm called **expectation propagation (EP)**
 - It requires us to be able to take expectations with respect to $p(\mathbf{z}|\mathbf{x})$!
 - In general, it's more computationally expensive than variational inference
 - Watch http://videolectures.net/mlss09uk_minka_ai/?q=minka

Variational inference (VI) vs expectation propagation (EP)

- Different KL leads to different properties for the approximation
 - $\mathbb{KL}(p(\mathbf{z}|\mathbf{x})||q(\mathbf{z}))$ in EP is moment-matching
 - $\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x}))$ in VI is mode-seeking

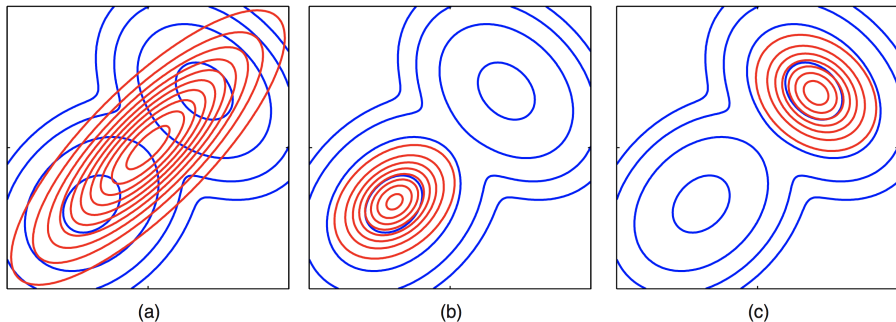


Figure: (a) $\mathbb{KL}(p||q)$. (b) $\mathbb{KL}(q||p)$. (c) Same as (b) but another local minimum. From Bishop (2006).

KL minimization

- Our goal is to find the variational parameters ν^* , such that

$$\nu^* = \arg \min_{\nu} \mathbb{KL}(q(\mathbf{z}|\nu) || p(\mathbf{z}|\mathbf{x})) = \arg \min_{\nu} \int_{\mathbf{z}} q(\mathbf{z}|\nu) \log \frac{q(\mathbf{z}|\nu)}{p(\mathbf{z}|\mathbf{x})}$$

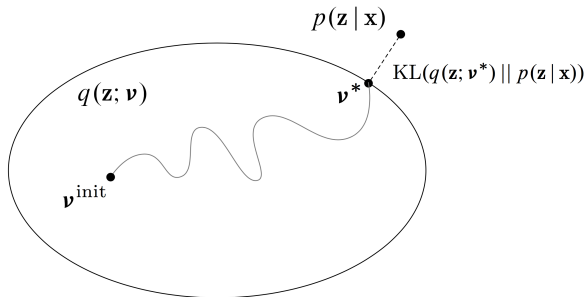


Figure: From David Blei (2017)

KL minimization

- Our goal is to find the variational parameters ν^* , such that

$$\nu^* = \arg \min_{\nu} \mathbb{KL}(q(\mathbf{z}|\nu) || p(\mathbf{z}|\mathbf{x})) = \arg \min_{\nu} \int_{\mathbf{z}} q(\mathbf{z}|\nu) \log \frac{q(\mathbf{z}|\nu)}{p(\mathbf{z}|\mathbf{x})}$$

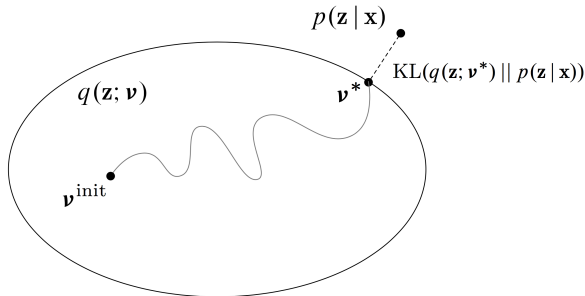


Figure: From David Blei (2017)

- Unfortunately, $\mathbb{KL}(q(\mathbf{z}) || p(\mathbf{z}|\mathbf{x}))$ cannot be minimized directly!

KL minimization

- However, we can find a function that we can minimize, which is equal to $\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x}))$ up to an additive constant, as follows

$$\begin{aligned}\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) &= \mathbb{E}_q \left[\log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right] \\ &= \mathbb{E}_q[\log q(\mathbf{z})] - \mathbb{E}_q[\log p(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_q[\log q(\mathbf{z})] - \mathbb{E}_q \left[\log \frac{p(\mathbf{z}, \mathbf{x})}{p(\mathbf{x})} \right] \\ &= - \underbrace{(\mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\mathbf{z})])}_{\mathcal{L}(q)} + \underbrace{\log p(\mathbf{x})}_{\text{const.}}\end{aligned}$$

KL minimization

- However, we can find a function that we can minimize, which is equal to $\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x}))$ up to an additive constant, as follows

$$\begin{aligned}\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) &= \mathbb{E}_q \left[\log \frac{q(\mathbf{z})}{p(\mathbf{z}|\mathbf{x})} \right] \\ &= \mathbb{E}_q[\log q(\mathbf{z})] - \mathbb{E}_q[\log p(\mathbf{z}|\mathbf{x})] \\ &= \mathbb{E}_q[\log q(\mathbf{z})] - \mathbb{E}_q \left[\log \frac{p(\mathbf{z}, \mathbf{x})}{p(\mathbf{x})} \right] \\ &= - \underbrace{(\mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\mathbf{z})])}_{\mathcal{L}(q)} + \underbrace{\log p(\mathbf{x})}_{\text{const.}}\end{aligned}$$

- The $\log p(\mathbf{x})$ term does not depend on q and thus it can be ignored
- Minimizing the KL divergence is then equivalent to maximizing $\mathcal{L}(q)$
- $\mathcal{L}(q)$ is called the **evidence lower bound (ELBO)**

KL minimization

- The ELBO, $\mathcal{L}(q)$, is a **lower bound** on the log model evidence $\log p(\mathbf{x})$

$$\mathbb{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x})) = \underbrace{-(\mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\mathbf{z})])}_{\text{ELBO } \mathcal{L}(q)} + \underbrace{\log p(\mathbf{x})}_{\text{log evidence}}$$

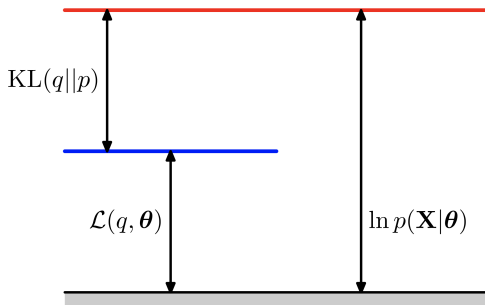


Figure: From Bishop (2006)

- The ELBO $\mathcal{L}(q)$ is tight when $q(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$, in which case $\mathcal{L}(q) \approx \log p(\mathbf{x})$

KL minimization

- The ELBO, $\mathcal{L}(q)$, is a **lower bound** on the log model evidence $\log p(\mathbf{x})$

$$\begin{aligned}\log p(\mathbf{x}) &= \log \int_{\mathbf{z}} p(\mathbf{z}, \mathbf{x}) \\ &= \log \int_{\mathbf{z}} q(\mathbf{z}) \frac{p(\mathbf{z}, \mathbf{x})}{q(\mathbf{z})} \\ &= \log \mathbb{E}_q \left[\frac{p(\mathbf{z}, \mathbf{x})}{q(\mathbf{z})} \right] \\ &\geq \underbrace{\mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\mathbf{z})]}_{\mathcal{L}(q)}\end{aligned}$$

- The last step comes from **Jensen's inequality**

$$\log \mathbb{E}[p(\mathbf{x})] \geq \mathbb{E}[\log p(\mathbf{x})]$$

- This is a consequence of the concavity of the logarithmic function

Variational inference

- Our goal is to find the variational parameters $\boldsymbol{\nu}$ that maximize the ELBO $\mathcal{L}(q)$

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\mathbf{z}|\boldsymbol{\nu})]$$

Variational inference

- Our goal is to find the variational parameters ν that maximize the ELBO $\mathcal{L}(q)$

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\mathbf{z}|\nu)]$$

- Can be maximized using a **coordinate ascent algorithm**
 - Iteratively optimize the variational parameters of each latent variable $q(z_m)$ in turn, holding the others fixed, until a convergence criteria is met
- Ensures convergence to a local maximum of $\mathcal{L}(q)$
- Remember, at convergence:

$$q(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$$

$$\mathcal{L}(q) \approx \log p(\mathbf{x})$$

Note

Variational inference (VI) is also often called **Variational Bayes (VB)**.

Two different perspectives on the ELBO $\mathcal{L}(q)$

- Expected value of log joint probability minus entropy terms

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \underbrace{\mathbb{E}_q[\log q(\mathbf{z}|\boldsymbol{\nu})]}_{\text{entropy terms}}$$

- Entropy terms favour less informative variational distributions $q(\mathbf{z}|\boldsymbol{\nu})$

Two different perspectives on the ELBO $\mathcal{L}(q)$

- Expected value of log joint probability minus entropy terms

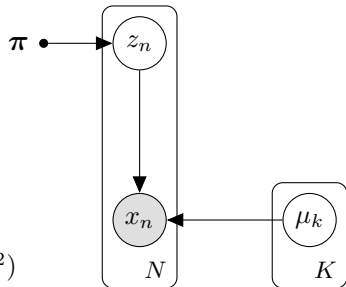
$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\mathbf{z}, \mathbf{x})] - \underbrace{\mathbb{E}_q[\log q(\mathbf{z}|\boldsymbol{\nu})]}_{\text{entropy terms}}$$

- Entropy terms favour less informative variational distributions $q(\mathbf{z}|\boldsymbol{\nu})$
- Expected value of log likelihood minus KL divergence to prior

$$\begin{aligned}\mathcal{L}(q) &= \mathbb{E}_q[\log p(\mathbf{x}|\mathbf{z})] - \mathbb{E}_q[\log p(\mathbf{z})] - \mathbb{E}_q[\log q(\mathbf{z}|\boldsymbol{\nu})] \\ &= \mathbb{E}_q[\log p(\mathbf{x}|\mathbf{z})] - \underbrace{\mathbb{E}_q\left[\log \frac{q(\mathbf{z}|\boldsymbol{\nu})}{p(\mathbf{z})}\right]}_{\mathbb{KL}(q(\mathbf{z}|\boldsymbol{\nu})||p(\mathbf{z}))}\end{aligned}$$

- $\mathbb{E}_q[\log p(\mathbf{x}|\mathbf{z})]$ is the reconstruction loss
- $\mathbb{KL}(q(\mathbf{z}|\boldsymbol{\nu})||p(\mathbf{z}))$ acts like a regularization term (to stay close to the prior)

A practical example: Bayesian Gaussian mixture model



- For the sake of simplicity, assume π is given (fixed)

① For each cluster k :

a) Draw cluster center $\mu_k \sim \mathcal{N}(\mu_k|0, \tau^2)$

② For each data point $1, \dots, N$:

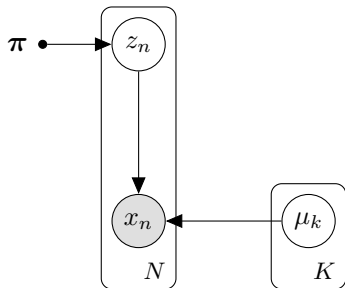
a) Draw cluster assignment $z_n \sim \text{Multinomial}(z_n|\pi)$

b) Draw observed data point $x_n \sim \mathcal{N}(x_n|\mu_{z_n}, \sigma^2)$

Note

For simplicity, we assumed the cluster variances σ^2 to be fixed.

A practical example: Bayesian Gaussian mixture model



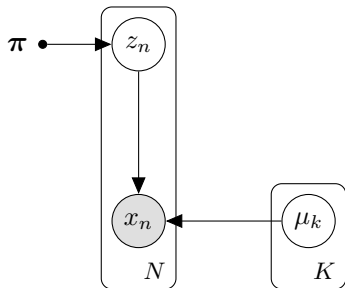
- Joint distribution:

$$p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x}) = \left(\prod_{k=1}^K p(\mu_k | 0, \tau^2) \right) \prod_{n=1}^N p(z_n | \boldsymbol{\pi}) p(x_n | z_n, \boldsymbol{\mu}, \sigma^2)$$

- Approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) =$$

A practical example: Bayesian Gaussian mixture model



- Joint distribution:

$$p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x}) = \left(\prod_{k=1}^K p(\mu_k | 0, \tau^2) \right) \prod_{n=1}^N p(z_n | \boldsymbol{\pi}) p(x_n | z_n, \boldsymbol{\mu}, \sigma^2)$$

- Approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K q(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N q(z_n | \phi_n)$$

A practical example: Bayesian Gaussian mixture model

- Our goal is to find the variational parameters that maximize the ELBO $\mathcal{L}(q)$

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\boldsymbol{\mu}, \mathbf{z})]$$

A practical example: Bayesian Gaussian mixture model

- Our goal is to find the variational parameters that maximize the ELBO $\mathcal{L}(q)$

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\boldsymbol{\mu}, \mathbf{z})]$$

- Because q is fully factorized, the expectations in $\mathcal{L}(q)$ decompose into sums of simpler terms

$$\begin{aligned}\mathbb{E}_q[\log p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})] &= \sum_{k=1}^K \mathbb{E}_q[\log p(\mu_k | 0, \tau^2)] + \sum_{n=1}^N \mathbb{E}_q[\log p(z_n | \boldsymbol{\pi})] \\ &\quad + \sum_{n=1}^N \mathbb{E}_q[\log p(x_n | z_n, \boldsymbol{\mu}, \sigma^2)]\end{aligned}$$

$$\mathbb{E}_q[\log q(\boldsymbol{\mu}, \mathbf{z})] = \sum_{k=1}^K \mathbb{E}_q[\log q(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k)] + \sum_{n=1}^N \mathbb{E}_q[\log q(z_n | \boldsymbol{\phi}_n)]$$

Expectations, expectations, expectations...

- For a discrete variable X :

$$\mathbb{E}[f(X)] = \sum_X p(X) f(X)$$

- For a continuous variable X :

$$\mathbb{E}[f(X)] = \int p(X) f(X) dX$$

- Some useful properties of expectations:

$$\mathbb{E}[a] = a$$

$$\mathbb{E}[a + bX] = a + b \mathbb{E}[X]$$

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$$

$$\mathbb{E}[XY] = \mathbb{E}[X] \mathbb{E}[Y], \quad \text{only if } X \text{ and } Y \text{ are independent}$$

$$\mathbb{V}[X] = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$$

- More details on expectations:

<https://www3.nd.edu/~rwilliam/stats1/x12.pdf>

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \phi_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] =$$

$$\mathbb{E}_q[\mu_k^2] =$$

$$\mathbb{E}_q[z_n] =$$

$$\mathbb{E}_q[z_{n,k}] =$$

$$\mathbb{E}_q[\mu_{z_n}] =$$

$$\mathbb{E}_q[\mu_{z_n}^2] =$$

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \boldsymbol{\phi}_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] = \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_k^2] =$$

$$\mathbb{E}_q[z_n] =$$

$$\mathbb{E}_q[z_{n,k}] =$$

$$\mathbb{E}_q[\mu_{z_n}] =$$

$$\mathbb{E}_q[\mu_{z_n}^2] =$$

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \boldsymbol{\phi}_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] = \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_k^2] = \tilde{\mu}_k^2 + \tilde{\sigma}_k^2$$

$$\mathbb{E}_q[z_n] =$$

$$\mathbb{E}_q[z_{n,k}] =$$

$$\mathbb{E}_q[\mu_{z_n}] =$$

$$\mathbb{E}_q[\mu_{z_n}^2] =$$

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \phi_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] = \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_k^2] = \tilde{\mu}_k^2 + \tilde{\sigma}_k^2$$

$$\mathbb{E}_q[z_n] = \sum_{k=1}^K k p(z_n = k) = \sum_{k=1}^K k \phi_{n,k}$$

$$\mathbb{E}_q[z_{n,k}] =$$

$$\mathbb{E}_q[\mu_{z_n}] =$$

$$\mathbb{E}_q[\mu_{z_n}^2] =$$

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \phi_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] = \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_k^2] = \tilde{\mu}_k^2 + \tilde{\sigma}_k^2$$

$$\mathbb{E}_q[z_n] = \sum_{k=1}^K k p(z_n = k) = \sum_{k=1}^K k \phi_{n,k}$$

$$\mathbb{E}_q[z_{n,k}] = p(z_n = k) = \phi_{n,k}$$

$$\mathbb{E}_q[\mu_{z_n}] =$$

$$\mathbb{E}_q[\mu_{z_n}^2] =$$

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \phi_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] = \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_k^2] = \tilde{\mu}_k^2 + \tilde{\sigma}_k^2$$

$$\mathbb{E}_q[z_n] = \sum_{k=1}^K k p(z_n = k) = \sum_{k=1}^K k \phi_{n,k}$$

$$\mathbb{E}_q[z_{n,k}] = p(z_n = k) = \phi_{n,k}$$

$$\mathbb{E}_q[\mu_{z_n}] = \sum_{k=1}^K \mathbb{E}_q[z_{n,k}] \mathbb{E}_q[\mu_k] = \sum_{k=1}^K \phi_{n,k} \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_{z_n}^2] =$$

Expectations, expectations, expectations...

- Consider the approximate (mean-field) distribution:

$$q(\boldsymbol{\mu}, \mathbf{z}) = \left(\prod_{k=1}^K \mathcal{N}(\mu_k | \tilde{\mu}_k, \tilde{\sigma}_k) \right) \prod_{n=1}^N \text{Mult}(z_n | \phi_n)$$

- Using the properties of expectations and the factorization of $q(\boldsymbol{\mu}, \mathbf{z})$:

$$\mathbb{E}_q[\mu_k] = \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_k^2] = \tilde{\mu}_k^2 + \tilde{\sigma}_k^2$$

$$\mathbb{E}_q[z_n] = \sum_{k=1}^K k p(z_n = k) = \sum_{k=1}^K k \phi_{n,k}$$

$$\mathbb{E}_q[z_{n,k}] = p(z_n = k) = \phi_{n,k}$$

$$\mathbb{E}_q[\mu_{z_n}] = \sum_{k=1}^K \mathbb{E}_q[z_{n,k}] \mathbb{E}_q[\mu_k] = \sum_{k=1}^K \phi_{n,k} \tilde{\mu}_k$$

$$\mathbb{E}_q[\mu_{z_n}^2] = \sum_{k=1}^K \mathbb{E}_q[z_{n,k}] \mathbb{E}_q[\mu_k^2] = \sum_{k=1}^K \phi_{n,k} (\tilde{\mu}_k^2 + \tilde{\sigma}_k^2)$$

Playtime!

- **Calculate** the coordinate ascent updates for maximizing $\mathcal{L}(q)$

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\boldsymbol{\mu}, \mathbf{z}, \mathbf{x})] - \mathbb{E}_q[\log q(\boldsymbol{\mu}, \mathbf{z})]$$

- Hints:

- Consider each variational parameter in turn: $\tilde{\mu}_k$, $\tilde{\sigma}_k$ and $\phi_{n,k}$

Table: Variables in the model and corresponding variational parameters

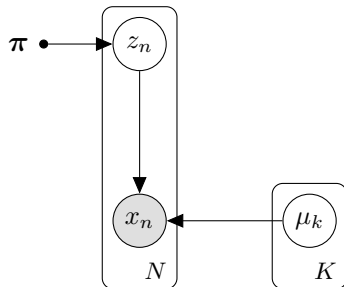
Model variable	μ_k		z_n
Variational parameter	$\tilde{\mu}_k$	$\tilde{\sigma}_k$	ϕ_n

- For each variational parameter, keep the relevant terms in $\mathcal{L}(q)$
- Take derivative with respect to the variational parameter and set it to zero
- If necessary, use Lagrange multipliers to ensure constraints
(e.g. $\sum_{k=1}^K \phi_{n,k} = 1$)

Note

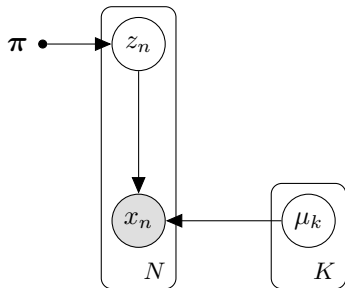
Automatic differentiation variational inference (ADVI) in STAN attempts to do all of this automatically (sometimes successfully, other times not so much).

Optimizing hyper-parameters



- What if we wanted to optimize the hyper-parameters π too?

Optimizing hyper-parameters



- What if we wanted to optimize the hyper-parameters π too?
- ELBO $\mathcal{L}(q)$ is a lower bound on the log marginal likelihood of the data $\log p(\mathbf{x})$

$$\log p(\mathbf{x}) = \log \int_{\boldsymbol{\mu}} \left(\prod_{k=1}^K p(\mu_k) \right) \prod_{n=1}^N \sum_{z_n} p(z_n | \pi) p(x_n | z_n, \boldsymbol{\mu}) \geq \mathcal{L}(q)$$

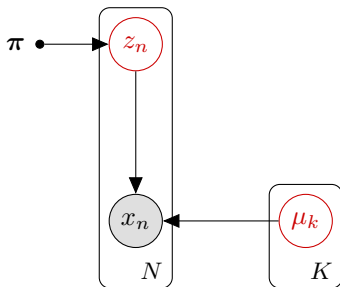
- At convergence of VI, this bound is tight: $\mathcal{L}(q) \approx \log p(\mathbf{x})$
- We can use $\mathcal{L}(q)$ to find a maximum likelihood estimate of π !

Variational Bayesian Expectation-Maximization (VB-EM)

- Goal is to find posterior over latent variables \mathbf{z} using VI and maximum likelihood estimates for hyper-parameters θ

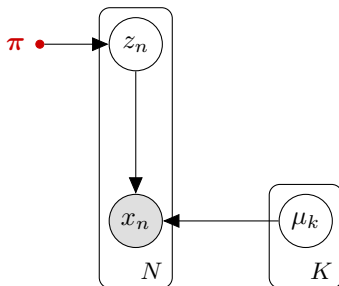
Variational Bayesian Expectation-Maximization (VB-EM)

- Goal is to find posterior over latent variables \mathbf{z} using VI and maximum likelihood estimates for hyper-parameters θ
- VB-EM alternates between these 2 steps until convergence:
 - E-step: run VI to find approximate posterior $q(\mathbf{z})$, such that $q(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$
 - M-step: use $\mathcal{L}(q) \approx \log p(\mathbf{x})$ to find maximum likelihood estimates of θ



Variational Bayesian Expectation-Maximization (VB-EM)

- Goal is to find posterior over latent variables \mathbf{z} using VI and maximum likelihood estimates for hyper-parameters θ
- VB-EM alternates between these 2 steps until convergence:
 - E-step: run VI to find approximate posterior $q(\mathbf{z})$, such that $q(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$
 - M-step: use $\mathcal{L}(q) \approx \log p(\mathbf{x})$ to find maximum likelihood estimates of θ



Relation with standard Expectation-Maximization (EM)

- Goal is to find posterior over latent variables \mathbf{z} using VI and maximum likelihood estimates for hyper-parameters θ
- VB-EM alternates between these 2 steps until convergence:
 - E-step: run VI to find approximate posterior $q(\mathbf{z})$, such that $q(\mathbf{z}) \approx p(\mathbf{z}|\mathbf{x})$
 - M-step: use $\mathcal{L}(q) \approx \log p(\mathbf{x})$ to find maximum likelihood estimates of θ

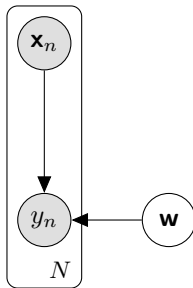
Note

If you are familiar with the popular Expectation-Maximization (EM) algorithm, it is a special case of VB-EM, when the E-step is exact (exact inference instead of VI).

VI example: Automatic Relevance Determination (ARD)

- Recall the Bayesian linear regression model:

- 1 Draw weights $\mathbf{w} \sim \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$
- 2 For each data point $n \in \{1, \dots, N\}$
 - a Draw target $y_n \sim \mathcal{N}(y_n|\mathbf{w}^T \mathbf{x}_n, \sigma^2)$

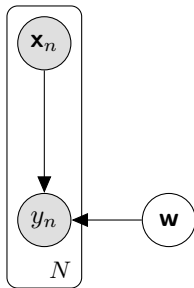


- It assumes independent priors on the weights \mathbf{w} with the same variance (α^{-1}) for all input dimensions $d \in \{1, \dots, D\}$

VI example: Automatic Relevance Determination (ARD)

- Recall the Bayesian linear regression model:

- 1 Draw weights $\mathbf{w} \sim \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$
- 2 For each data point $n \in \{1, \dots, N\}$
 - a Draw target $y_n \sim \mathcal{N}(y_n|\mathbf{w}^T \mathbf{x}_n, \sigma^2)$



- It assumes independent priors on the weights \mathbf{w} with the same variance (α^{-1}) for all input dimensions $d \in \{1, \dots, D\}$
- Alternatively, we consider a different variance α_d^{-1} for each input dimension
- We can then specify a prior on the precisions $\alpha_d \sim \text{Gamma}(\alpha_d|a_0, b_0)$
- This is called an **hyper-prior**!

VI example: Automatic Relevance Determination (ARD)

- New generative process and PGM:

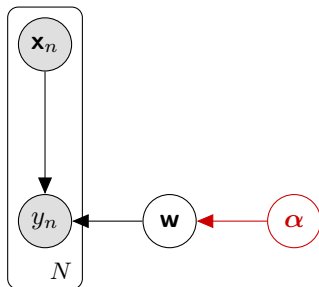
① For each input dimension $d \in \{1, \dots, D\}$

 a Draw weight precision
 $\alpha_d \sim \text{Gamma}(\alpha_d | a_0, b_0)$

② Draw weights $\mathbf{w} \sim \mathcal{N}(\mathbf{w} | \mathbf{0}, \alpha^{-1} \mathbf{I})$

③ For each data point $n \in \{1, \dots, N\}$

 a Draw target $y_n \sim \mathcal{N}(y_n | \mathbf{w}^T \mathbf{x}_n, \sigma^2)$



- Prior on α_d specifies how uncertain we are a-priori that these weights are small

VI example: Automatic Relevance Determination (ARD)

- New generative process and PGM:

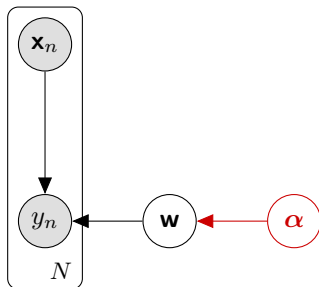
① For each input dimension $d \in \{1, \dots, D\}$

 a Draw weight precision
 $\alpha_d \sim \text{Gamma}(\alpha_d | a_0, b_0)$

② Draw weights $\mathbf{w} \sim \mathcal{N}(\mathbf{w} | \mathbf{0}, \alpha^{-1} \mathbf{I})$

③ For each data point $n \in \{1, \dots, N\}$

 a Draw target $y_n \sim \mathcal{N}(y_n | \mathbf{w}^T \mathbf{x}_n, \sigma^2)$



- Prior on α_d specifies how uncertain we are a-priori that these weights are small
- Posterior on α_d tells us how relevant the d^{th} dimension is - hence the name ARD!
- Added flexibility allows some weights w_d to be further pushed towards zero
- It allows to do **automatic feature selection**!

VI example: Automatic Relevance Determination (ARD)

- The joint distribution is given by:

$$p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \boldsymbol{\alpha}^{-1} \mathbf{I}) \left(\prod_{j=1}^D \text{Ga}(\alpha_j | a_0, b_0) \right) \prod_{n=1}^N \mathcal{N}(y_n | \mathbf{w}^T \mathbf{x}_n, \sigma^2)$$

where $\mathbf{y} = \{y_1, \dots, y_N\}$ and $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$

VI example: Automatic Relevance Determination (ARD)

- The joint distribution is given by:

$$p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \boldsymbol{\alpha}^{-1} \mathbf{I}) \left(\prod_{j=1}^D \text{Ga}(\alpha_j | a_0, b_0) \right) \prod_{n=1}^N \mathcal{N}(y_n | \mathbf{w}^T \mathbf{x}_n, \sigma^2)$$

where $\mathbf{y} = \{y_1, \dots, y_N\}$ and $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$

- Exact inference is unfortunately now intractable

$$p(\mathbf{w}, \boldsymbol{\alpha} | \mathbf{y}, \mathbf{X}) = \frac{p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X})}{\int_{\mathbf{w}} \int_{\boldsymbol{\alpha}} p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X})}$$

VI example: Automatic Relevance Determination (ARD)

- The joint distribution is given by:

$$p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \boldsymbol{\alpha}^{-1} \mathbf{I}) \left(\prod_{j=1}^D \text{Ga}(\alpha_j | a_0, b_0) \right) \prod_{n=1}^N \mathcal{N}(y_n | \mathbf{w}^T \mathbf{x}_n, \sigma^2)$$

where $\mathbf{y} = \{y_1, \dots, y_N\}$ and $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$

- Exact inference is unfortunately now intractable

$$p(\mathbf{w}, \boldsymbol{\alpha} | \mathbf{y}, \mathbf{X}) = \frac{p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X})}{\int_{\mathbf{w}} \int_{\boldsymbol{\alpha}} p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X})}$$

- But we can use (mean-field) variational inference!

$$q(\mathbf{w}, \boldsymbol{\alpha}) = \mathcal{N}(\mathbf{w} | \mathbf{m}_N, \mathbf{V}_N) \prod_{j=1}^D \text{Ga}(\alpha_j | a_{Nj}, b_{Nj})$$

- Our goal is to find the variational parameters \mathbf{m}_N , \mathbf{V}_N , a_{Nj} and b_{Nj} that maximize the ELBO $\mathcal{L}(q)$:

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X})] - \mathbb{E}_q[\log q(\mathbf{w}, \boldsymbol{\alpha})]$$

Playtime!

- **Implement** the coordinate ascent updates for maximizing $\mathcal{L}(q)$

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(\mathbf{y}, \mathbf{w}, \boldsymbol{\alpha} | \mathbf{X})] - \mathbb{E}_q[\log q(\mathbf{w}, \boldsymbol{\alpha})]$$

- Jupyter notebook: "12 - Variational inference - ARD.ipynb"
- Hints:
 - Consider each variational parameter in turn: \mathbf{m}_N , \mathbf{V}_N , a_{Nj} and b_{Nj}

Table: Variables in the model and corresponding variational parameters

Model variable	\mathbf{w}		α_j	
Variational parameter	\mathbf{m}_N	\mathbf{V}_N	a_{Nj}	b_{Nj}

Note

Notice how fast this VI algorithm is when compared to MCMC inference in STAN!